Project Red Goose

Neural Interdiction and Remote Signal Research Framework
Revised publication-style manuscript with grounded clinical framing and quarantined exploratory content

Charles “Eyes” Williams | Smalleyes 0m364
April 19, 2026

Document purpose	Reframe the original concept paper into a cleaner research manuscript by separating clinically grounded architecture from speculative long-range sensing theory.
Primary contribution	A deterministic, interrupt-driven orchestration model for multimodal seizure detection and rapid neuromodulatory response, with generative AI restricted to interpretation-only pathways.
Editorial note	Phase 2 remains hypothesis-generating only and is presented as exploratory physics-facing research rather than a deployable capability.

Prepared from the user-supplied manuscript and editorially revised for clarity, structure, and technical defensibility.

Abstract

Project Red Goose is presented here as a two-tier research program. Tier 1 is a clinically grounded concept for closed-loop seizure interdiction in refractory post-traumatic epilepsy. It combines multimodal sensing, deterministic signal classification, and non-ablative ultrasound neuromodulation under a hard real-time orchestration model intended to minimize response latency. Tier 2 is a deliberately segregated exploratory research track concerning whether very weak neural-field information might ever be characterized at distance under tightly controlled, fully consented conditions. The central engineering claim of the paper is not that speculative remote sensing is presently achievable, but that any credible future investigation would first require a rigorously partitioned Tier 1 platform in which signal acquisition, classification, and actuation are auditable and latency-bounded. The manuscript argues for three design principles: first, generative AI must be excluded from execution pathways; second, multimodal concurrence should be required before therapeutic action; and third, exploratory remote-acquisition theory must be treated as hypothesis generation, not capability assertion. Recent literature on neuromorphic neuromodulation, hybrid EEG-fNIRS systems, and focused ultrasound supports the architectural plausibility of the Tier 1 stack, while the Tier 2 discussion remains explicitly constrained by major physical, ethical, and noise-floor barriers.

Keywords: closed-loop neuromodulation; seizure interdiction; traumatic brain injury epilepsy; EEG; fNIRS; focused ultrasound; RFML; deterministic orchestration; neuromorphic classification; exploratory remote neural sensing

Editorial Revisions Introduced in This Version

- Phase 1 was rewritten as the primary scientific contribution, with publication-style tone and tighter claims.
- Phase 2 was moved into a clearly bounded exploratory section and stripped of any implication of present operational feasibility.
- The execution-layer argument was sharpened: language models may summarize logs and assist operators, but do not touch the raw signal path or stimulation path.
- Risk language was expanded to include false positives, adversarial injection, synchronization drift, and ethical governance limits.
- A short falsifiability framework was added so the speculative portion reads like a testable research question instead of a cinematic mood board.

1. Introduction

Traumatic brain injury is a major cause of acquired epilepsy, and a substantial fraction of patients with epilepsy do not achieve seizure freedom with medication alone. For refractory cases, currently deployed options such as resection, vagus nerve stimulation, deep brain stimulation, and responsive neurostimulation can be clinically valuable but remain constrained by invasiveness, access, targeting tradeoffs, and system latency.

Project Red Goose begins from a narrower and more defensible premise than the original draft occasionally implied. Its core problem is not mind reading, intention inference, or generalized neural decoding at distance. Its core problem is whether an engineered, multimodal, hard-partitioned control stack could detect pre-ictal transition markers fast enough to trigger a reversible suppressive intervention before full seizure escalation in selected patients.

The paper therefore treats Red Goose as an orchestration problem. The issue is not merely which sensors or classifiers to use, but how to assign responsibility across layers so that the system remains auditable, latency-bounded, and clinically interpretable. The proposed architecture isolates natural-language systems from the therapeutic pathway and reserves high-speed decisions for deterministic components whose behavior can be traced and tested.

2. Technical Background

Three lines of literature matter most to the present framework. The first is closed-loop neuromodulation and the long-standing problem of end-to-end latency. The second is multimodal non-invasive sensing, especially EEG paired with fNIRS, where one modality contributes temporal precision and the other contributes complementary physiological context. The third is focused ultrasound neuromodulation, which has moved from foundational preclinical work toward early human safety studies in epilepsy and related neurologic indications.

Recent perspective work on neuromorphic neuromodulation argues that low-power, on-device classifiers may be better suited than cloud-linked or telemetry-heavy pipelines for implantable or wearable closed-loop systems. That literature does not solve the seizure-interdiction problem by itself, but it supports the paper's insistence that classification must remain near the signal source rather than being pushed through variable-latency interpretive layers.

Hybrid EEG-fNIRS systems have likewise become more plausible as ambulatory research platforms. EEG supplies millisecond-scale temporal information and established seizure biomarker traditions, while fNIRS contributes hemodynamic context and can improve robustness when electrical signatures are ambiguous or contaminated by artifact. Focused

ultrasound, particularly low-intensity regimes, is relevant not because it magically fixes epilepsy, but because it offers a non-ablative route to targeted neuromodulation that can be investigated without collapsing the architecture into a purely implant-dependent paradigm.

3. Architecture of the Red Goose Tier 1 Platform

The proposed Tier 1 architecture is divided into three layers: Interface, Logic, and Orchestration. The Interface layer may contain a large language model or similar assistant, but only for operator dialogue, summarization, and post hoc interpretation. It receives processed logs, confidence summaries, and system status, not raw packets. The Logic layer performs deterministic feature extraction and state classification. The Orchestration layer is a real-time arbitration engine that allocates sensing and actuation resources, enforces concurrence requirements, and issues or withholds therapeutic commands.

This separation is the manuscript's principal engineering contribution. In many AI-adjacent medical concept papers, 'intelligence' is treated as a monolith. Here it is deliberately disassembled. Human-readable reasoning and natural-language convenience belong in one compartment. High-speed classification belongs in another. Hardware scheduling and actuation belong in a third. That partition is what allows auditability without sacrificing responsiveness.

The functional objective is simple to state and hard to achieve: detect a pre-ictal transition, validate it across modalities, and deliver a reversible suppressive pulse while false positives remain acceptably low. Any part of the stack that introduces stochastic, variable-latency behavior into that path is therefore suspect by default.

3.1 Interface Layer

The Interface layer exists for operators, researchers, and patients. It can summarize trends, explain why a trigger was or was not issued, display recent anomaly scores, and translate internal state into human-readable language. It may support maintenance tasks, threshold review, or anomaly-logging workflows.

It may not classify raw neural packets in the live control loop. It may not synthesize therapeutic actions. It may not directly schedule stimulation. Restricting it this way is not anti-AI theatrics. It is the only sane response to the fact that probabilistic language systems are variable in both output and timing.

3.2 Logic Layer

The Logic layer is framed using RFML-adjacent thinking: the incoming stream is treated as a waveform environment in which relevant state transitions can be learned or thresholded from amplitude, phase-related, spectral, and distortion features. The point is not that the brain is literally a radio. The point is that signal-intelligence methods are often better conceptual tools for pre-ictal transition detection than vague talk about 'reading thoughts.'

In this manuscript, the classifier is deterministic by design. A trainable model may be used offline to derive features or tune thresholds, but the live execution path should compile down to bounded, inspectable logic. That may include feature vectors, event windows, concurrence scoring, and rule-based gating. It should not include an unconstrained generative model making improvisational decisions in the middle of a medical event, because apparently someone somewhere always has to learn that lesson the expensive way.

3.3 Orchestration Layer

The Orchestration layer functions as an interrupt-driven real-time controller. When the Logic layer crosses a configured threshold, the Orchestration layer checks temporal alignment, modality concurrence, safety interlocks, refractory periods, and hardware readiness before issuing a suppression command. If concurrence fails, if synchronization is stale, or if the signal quality estimate drops below threshold, the action is blocked and the event is logged for review.

This layer also manages the shared use of hardware resources. If an ultrasound transducer or related subsystem must alternate between sensing-oriented and actuation-oriented roles, arbitration belongs here rather than in the classifier itself. Keeping resource policy separate from classification prevents hidden coupling and makes failure analysis tractable.

4. Sensor and Actuation Stack

Tier 1 assumes a non-invasive or minimally invasive research configuration that combines EEG, fNIRS, and focused ultrasound-related hardware. EEG remains the fastest practical non-invasive source for electrical state transitions, but it is noisy and vulnerable to motion artifact. fNIRS is slower because hemodynamic responses lag, yet it can still contribute valuable corroborative information across longer windows or recurring patient-specific pre-ictal patterns.

The manuscript's original triad is retained conceptually but tightened in language. The sensors are not mystical antennae and the actuation hardware is not a generic

'metronome.' They are simply components of a closed-loop system with different latencies, noise profiles, and safety envelopes. Their value comes from complementary failure modes, not from grandiosity.

Actuation is modeled as low-intensity focused ultrasound or a similarly reversible suppressive modality under strict safety constraints. The paper does not claim that universal sub-millisecond seizure termination has already been shown in humans. It argues instead that a deterministic multimodal controller paired with a reversible suppressive method is a more coherent experimental path than architectures that mix interpretation, telemetry, and actuation into one messy box.

5. Safety, Latency, and Adversarial Considerations

Latency is not just a performance metric. In this context it is a safety boundary. Excessive delay undermines utility, but poorly governed speed can also produce dangerous false positives. The architecture therefore uses concurrence, refractory logic, and hard interlocks rather than assuming that faster is always better.

The manuscript also retains the threat model of signal injection or spoofing, though in stricter language. Any external telemetry or maintenance channel creates opportunities for corruption, replay, or malformed input. For that reason, live therapeutic decisions should rely on local validated sensor pathways, authenticated timing sources, and encrypted communications for all external links. Cross-modal agreement is not merely a statistical convenience. It is a defensive design choice.

Another critical risk is synchronization drift. Multimodal systems often look elegant on whiteboards and become ridiculous once clocks diverge, sensors slip, or patient motion pushes one stream just far enough out of alignment to fake a pattern. Tier 1 must therefore include continuous alignment checks, known-latency calibration procedures, and conservative fallbacks.

6. Tier 2 Exploratory Research Track: Remote Neural Signal Characterization

The original paper's Phase 2 has been retained only because removing it entirely would erase the motivating frontier that inspired the project. It is now restated more precisely: Tier 2 asks whether extremely weak neural-field-associated information could ever be characterized at standoff distance under laboratory conditions using exotic sensing or field-coupling approaches. It does not assert that current systems can do this, and it does not describe a deployable surveillance capability.

The dominant barrier is still the same one physics keeps rudely enforcing: the neural signals observable at the scalp are already weak, and their electromagnetic or mechanical correlates attenuate rapidly with distance while drowning in environmental noise. Any

credible remote characterization program would therefore need to prove signal existence, coupling mechanism, and reproducibility under blinded controls before grander interpretations are even entertained.

The manuscript mentions theoretical work on weakly evanescent cortical waves and related field models only as motivation for asking whether standard quasi-static assumptions are complete in all contexts. That literature does not establish practical standoff brain reading. At best, it suggests that the brain's field behavior may be richer than oversimplified models imply.

6.1 On Sonoluminescence and Plasma-Assisted Hypotheses

The prior draft proposed a plasmatic enhancement hypothesis involving sonoluminescence-related phenomena. That idea is preserved here as a clearly labeled hypothesis-generating concept and nothing more. Sonoluminescence demonstrates that collapsing cavitation bubbles can create intense localized conditions, but translating that fact into a safe, selective, information-amplifying neural sensing mechanism is an enormous leap across several missing bridges.

The most immediate problems are obvious. Cavitation is spatially localized. Neural tissue safety margins are unforgiving. Environmental control is difficult. Any induced plasma-related effect would have to improve signal discrimination rather than merely adding spectacular noise. At present, no published evidence shows a practical route from sonoluminescent physics to non-contact neural-state acquisition in the manner imagined by the original concept.

That does not make the idea forbidden to discuss. It means the burden of proof is severe. In this manuscript, the hypothesis is retained only as an example of a falsifiable frontier question, not as a latent promise.

6.2 Falsifiability Framework for Tier 2

To keep the exploratory track scientific, the manuscript adopts a simple falsifiability ladder. First, demonstrate any reproducible above-noise remote correlate from a consented, instrumented subject under shielded laboratory conditions. Second, show that the correlate tracks controlled state transitions better than chance under blinding. Third, demonstrate that the proposed coupling mechanism contributes explanatory value beyond conventional artifact sources. Fourth, replicate across independent equipment and operators.

If those conditions fail, then Tier 2 remains an interesting idea and nothing more. Which, frankly, is still better than half the things people attach the word 'quantum' to in public.

7. Ethical and Governance Constraints

Tier 1 already raises serious ethical issues because a false positive or misconfigured suppressive intervention can directly affect a patient. Tier 2 raises an additional category of concerns around privacy, surveillance, and consent. For that reason, the manuscript adopts a bright line: any exploratory work on remote neural characterization is ethically discussable only in fully consented, instrumented, review-governed settings. No covert or non-consensual use case is endorsed or implied.

The governance model should also distinguish between detecting a physiological state transition and decoding semantic content. These are not the same problem. Conflating them is both technically sloppy and ethically dangerous. Any future architecture even approaching content reconstruction would require a substantially different legal and ethical framework.

8. Discussion

The revised manuscript is strongest where it remains architectural. The value of Red Goose is not that it magically solves epilepsy or establishes remote neural sensing. The value is that it proposes a disciplined systems pattern for multimodal detection and actuation under latency pressure. That pattern can be tested, falsified, refined, or rejected using normal engineering criteria.

Tier 1 aligns with active research trajectories: low-power edge classification, hybrid sensing, and focused ultrasound neuromodulation. None of those fields alone validates the full Red Goose stack, but together they make the platform plausible enough to justify staged prototyping and safety-first investigation. Tier 2, by contrast, remains a distant speculative track whose merit depends entirely on future empirical evidence.

That asymmetry is intentional. A serious paper does not gain credibility by pretending all of its ideas are equally mature. It gains credibility by admitting which parts stand on literature, which parts stand on analogy, and which parts are still standing out in the rain waiting for data.

9. Conclusion

Project Red Goose is best understood as a layered research framework rather than a single claim. Its grounded contribution is a partitioned control architecture for seizure interdiction research in which deterministic classification and real-time orchestration are kept separate from natural-language assistance. That architecture is compatible with the current direction of work in neuromorphic neuromodulation, hybrid EEG-fNIRS sensing, and focused ultrasound neuromodulation.

The manuscript's exploratory contribution is narrower than before and therefore stronger. It identifies remote neural signal characterization as a frontier question, specifies why present barriers remain formidable, and provides a falsifiability ladder rather than a theatrical promise. In other words, it now reads less like forbidden wizardry and more like an actual research program, which is generally an improvement.

Tier 1 Layer Summary

Layer	Primary Role	Permitted Data	Disallowed Function
Interface	Operator interaction, logs, status, explanation	Processed summaries and event logs	Live classification, raw neural stream access, direct actuation
Logic	Deterministic feature extraction and state classification	Validated sensor streams and derived features	Natural-language reasoning in live control loop
Orchestration	Real-time arbitration, interlocks, command gating	Threshold events, timing state, hardware readiness	Probabilistic interpretation or free-form decision making

Selected References

[1] Contreras LFH, Truong ND, Eshraghian JK, et al. Neuromorphic neuromodulation: Towards the next generation of closed-loop neurostimulation. PNAS Nexus. 2024;3(11):pgae488.

[2] Li B, Li M, Xia J, et al. Hybrid integrated wearable patch for brain EEG-fNIRS monitoring. Sensors. 2024;24(15):4847.

[3] Agopyan-Miu A, Simmons GB, Baltuch GH. Is there a role for focused ultrasound in the treatment of refractory epilepsy? Stereotactic and Functional Neurosurgery. 2025.

[4] Hogeveen L, et al. Therapeutic use of transcranial ultrasound for epilepsy. Heliyon. 2025.

[5] Min BK, Bystritsky A, Jung KI, et al. Focused ultrasound-mediated suppression of chemically induced acute epileptic EEG activity. BMC Neuroscience. 2011;12:23.

[6] West N, et al. An RFML ecosystem: Considerations for the application of deep learning to spectrum situational awareness. IEEE Open Journal of the Communications Society. 2021;2:2243-2264.

[7] Galinsky VL, Frank LR. Brain waves: Emergence of localized, persistent, weakly evanescent cortical wave loops. Journal of the Royal Society Interface. 2020.

[8] Szymanski RA, et al. Principles and applications of sono-optogenetics. 2023.

Appendix A. Suggested Next-Step Research Sequence

1. Bench validation of multimodal timing and alignment on phantom signals before any patient-facing work.
2. Offline feature study to identify patient-specific pre-ictal markers and estimate false-positive burden.
3. Closed-loop dry-run mode with actuation disabled to test orchestration and concurrence policy.
4. Safety-reviewed actuation studies limited to pre-defined windows and conservative parameters.
5. Only after Tier 1 instrumentation proves stable should any Tier 2 laboratory exploration be considered.